

# 104 - Easily switch infrastructure and manage teams with work pool-based deployments

# 104 Agenda

---

- Create work pool-based deployments with `.deploy()`
- Run flows on Prefect's infrastructure with a Prefect Managed work pool
- Use a worker with a hybrid work pool for maximum control
  - Process
  - Docker
- Store your flow code
  - On GitHub
  - In a Docker image



## Why use a work pool-based deployment?

---

Infrastructure is a pain, Prefect makes it better. 😊

- Run workflows on a variety of dynamic infrastructure
- Provide a template for teams
- Scale infrastructure to 0 (serverless)
- Prioritize work

## Create deployment with `.deploy()`

---


Very similar syntax to `.serve()`

Differences:

- Must specify work pool
- Must specify flow code storage source (or Docker image)
- Does **not** start a process watching for scheduled runs



# Managed work pools



# Prefect Managed work pools

---

- Run workflows on Prefect's infrastructure
- Cloud only
- Easy mode - no worker required
- Limitations
  - Compute hours
  - Concurrency
  - No custom Docker image



## First work pool-based deployment

---

- Create deployment with *.deploy()*
- Specify flow code stored in a GitHub repository with *.from\_source()*
- Use a **Prefect Managed** work pool



# Create a Prefect Managed work pool

In the UI, **Work Pools** -> + -> **Prefect Managed**

**Work Pools / Create**

✓ Infrastructure Type    02 Details    03 Configuration

Name

Description (Optional)

Flow Run Concurrency (Optional)

Unlimited

Cancel Previous Next





## Work pools

---

- Don't modify the job template for now
- You can specify environment variables, etc.
- Work pools make it easier for data engineering platform teams to create guardrails for other teams



# Deployment with Prefect Managed work pool

---

```
from prefect import flow

if __name__ == "__main__":
    flow.from_source(
        source="https://github.com/biancaines/pal-2025-v1.git",
        entrypoint="102/weather2-tasks.py:pipeline",
    ).deploy(
        name="my-first-managed-deployment",
        work_pool_name="my-managed-workpool",
    )
```



# Run script to create the deployment

---

Successfully created/updated all deployments!

## *Deployments*

Name	Status	Details
pipeline/my-first-managed-deployment	applied	

To schedule a run for this deployment, use the following command:

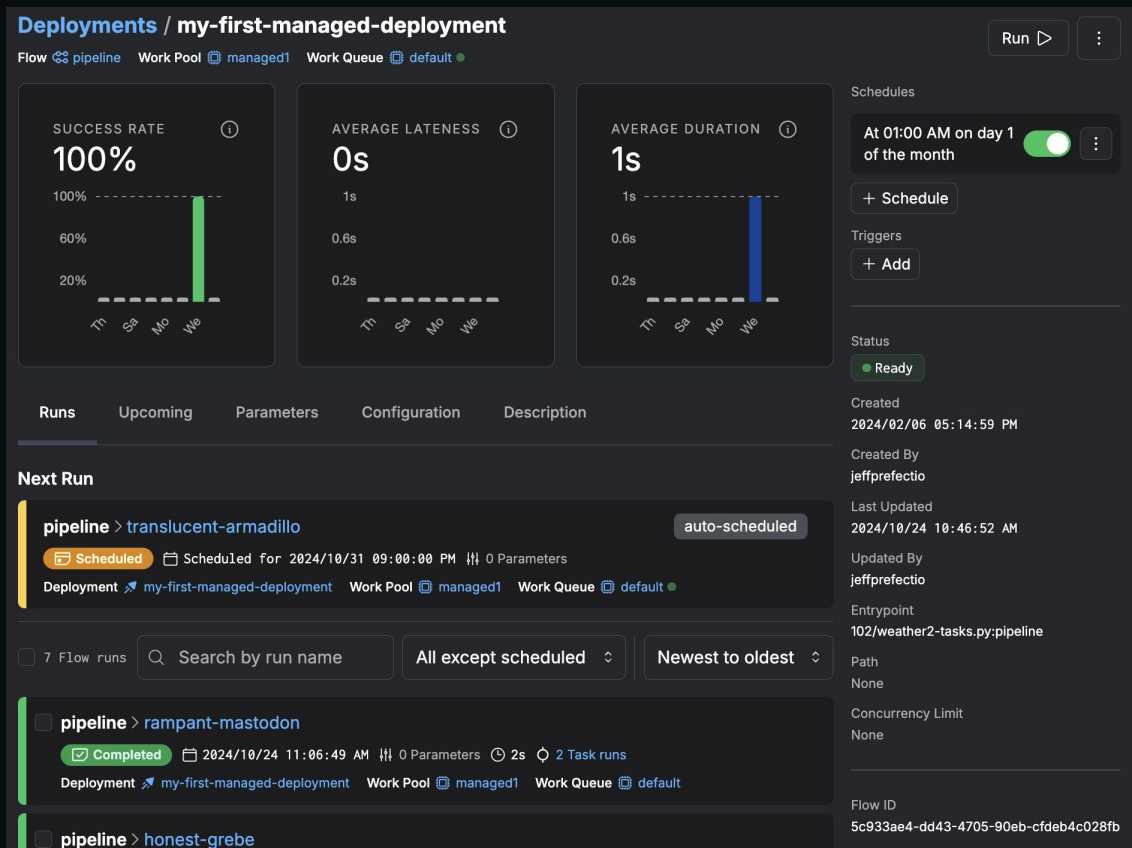
```
$ prefect deployment run 'pipeline/my-first-managed-deployment'
```

You can also run your flow via the Prefect UI: <https://app.prefect.cloud/account/9b649228-0419-40e1-9e0d-44954b5c0ab6/workspace/d137367a-5055-44ff-b91c-6f7366c9e4c4/deployments/deployment/d448be8f-2092-47f9-8d0b-ee06ce182480>



# See deployment details in the UI

1.



# Run the deployment

Runs / pipeline / my-first-managed-deployment / honest-grebe Completed

10:52:33 AM 10:52:34 AM 10:52:35

fetch\_weather-72f save\_weather-bba

Filter Search Display

honest-grebe +10 10:52:33 AM

- prefect.flow-run.Scheduled 10:52:00 AM
- prefect.flow-run.Pending 10:52:11 AM
- Opening process... 10:52:28 AM
- prefect.flow-run.Running 10:52:33 AM
  - fetch\_weather-72f +5 10:52:33 AM
    - save\_weather-bba +4 10:52:34 AM
      - Successfully wrote temp 10:52:34 AM
      - prefect.flow-run.Completed 10:52:34 AM
      - Finished in state Completed() 10:52:34 AM
      - Process for flow run 'honest-... 10:52:37 AM

honest-grebe Flow run

Automation None

Tags

Add tags to your task and flow definitions to categorize them and enable concurrency limiting. [Learn more](#)

Parameters

Add parameters to your flow definitions to make them reusable and configurable. [Learn more](#)

Configuration

Retries 3

Retry delay 0

Retry jitter None

Flow version 5773fc75feaf48a8a81f6f1850c7f32b

Logs

Level: all

Opening process... 10:52:28 AM INFO

Forecasted temp C: 18.6 degrees 10:52:34 AM INFO

Finished in state Completed() 10:52:34 AM INFO

Finished in state Completed() 10:52:34 AM INFO

Successfully wrote temp 10:52:34 AM INFO

Finished in state Completed() 10:52:34 AM INFO



## At runtime, Prefect:


---

1. Pulls your flow code from GitHub
2. Runs your code in a Docker container on our infrastructure
3. Monitors and reports on state
4. Exits container and cleans up



# Run the deployment

---

- Run state progression:  
*Scheduled -> Pending -> Running -> Completed*
-  Takes a moment to spin up Docker Container on our infrastructure



# Let's break this down

---





# Work pools



## Work pools

---

- Server side
- Provide default infrastructure configuration for deployments
- 🖐️ **Deployments that use this work pool inherit these settings**



# Flow code storage



# Flow code storage options


---

1. Local
2. Git-based remote repository (e.g. GitHub, GitLab)
3. Bake your code into a Docker image
4. Cloud provider storage



## Flow code storage

---

- Specified public GitHub repo with *.from\_source()* class method
- Call *flow.from\_source()* or *flow\_name.from\_source()*
- Provide repo URL and *entrypoint path:flow function name*
-  Private repos are fine, just pass credentials




# Hybrid model - hybrid work pools with workers



# Hybrid work pools with workers

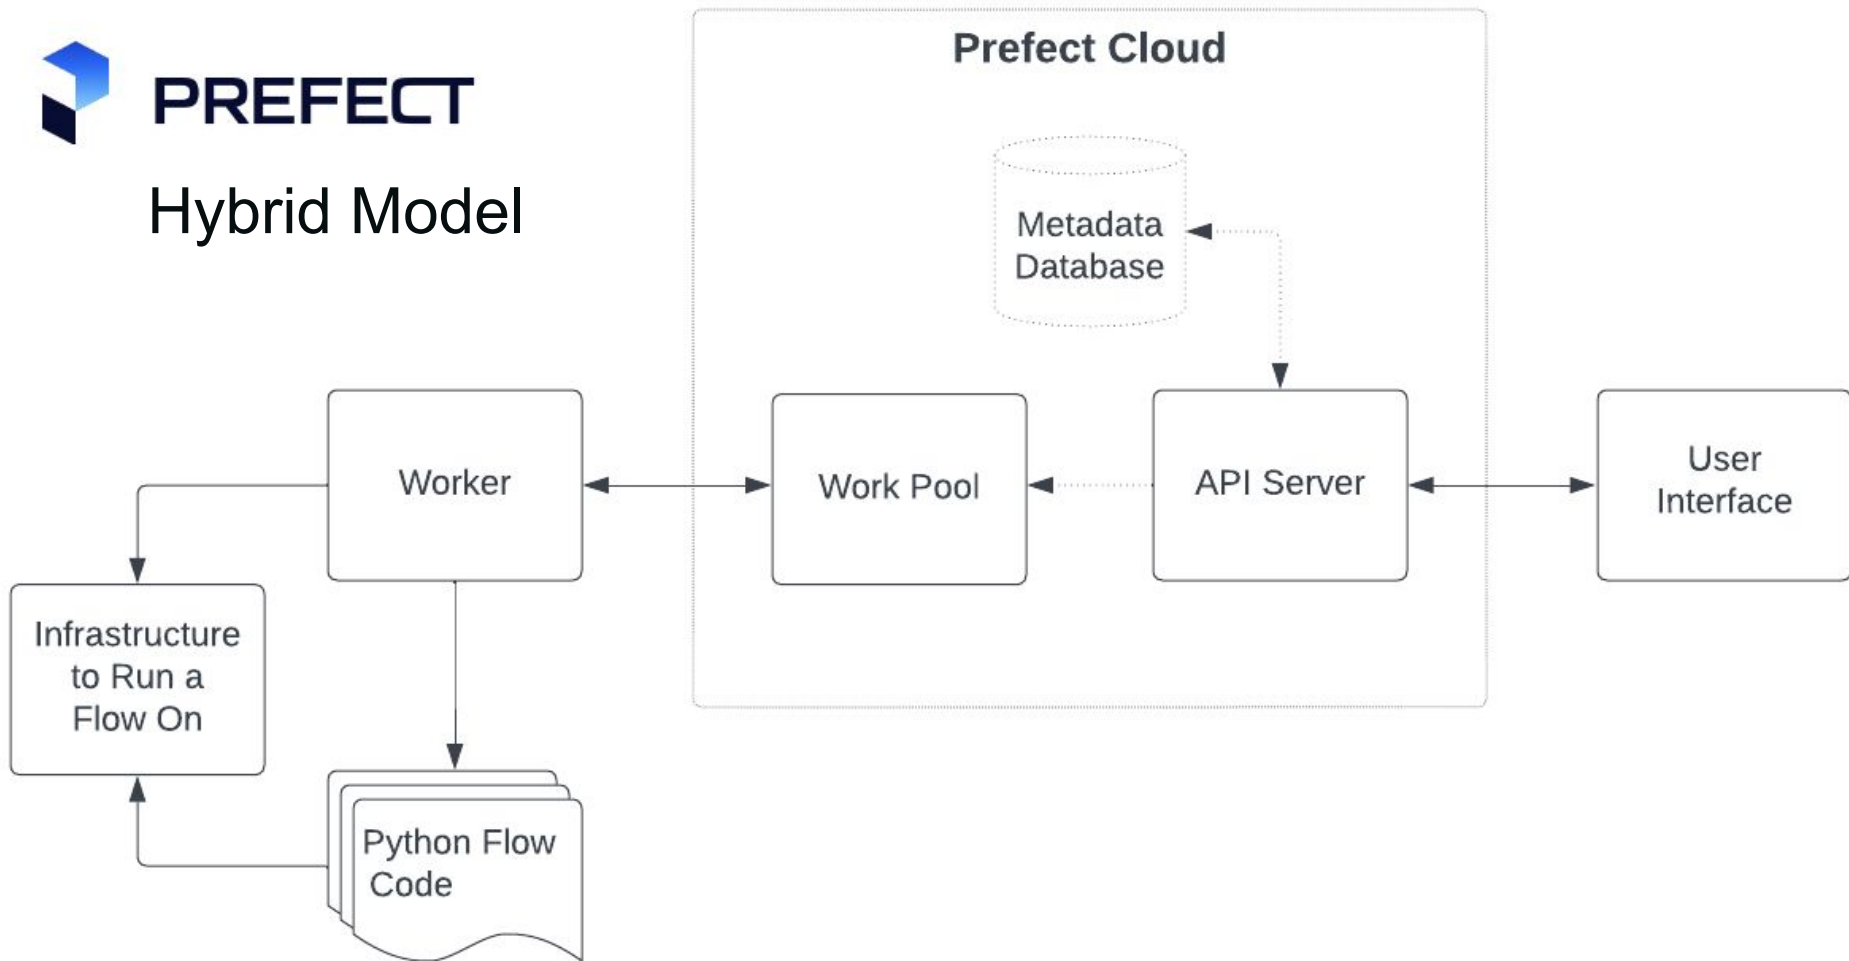
---





# PREFECT

## Hybrid Model





## Hybrid model = separation

---

- Your flow code runs on your infrastructure
- Your flow code is stored on your storage (GitLab, GitHub, AWS, Docker image, etc)
- Prefect Cloud stores metadata, logs, artifacts, etc.
- Data encrypted at rest
- Prefect Technologies, Inc. is SOC2 Type II compliant

<https://www.prefect.io/security>



# Example: Process work pool & worker



## First hybrid work pool-based deployment

---

- Create with `.deploy()`
- Specify flow code stored in a GitHub repository with `.from_source()`
- Use a **Process** work pool
- Start a worker to pick up scheduled flow runs ★



# Create a Process work pool

---



## Process

Execute flow runs as subprocesses on a worker. Works well for local execution when first getting started.



# Create a Process work pool

**Work Pools / Create**

✓ Infrastructure Type

✓ Details

03 Configuration

Below you can configure defaults for deployments that use this work pool. Use the editor in the **Advanced** section to modify the existing configuration options, if needed.

If you don't need to change the default configuration, click **Create** to create your work pool!

Base Job Template

DefaultsAdvanced

1

The fields below control the default values for the base job template. These values can be overridden by deployments.

**Name (Optional)**  
Name given to infrastructure created by a worker.

**Environment Variables (Optional)**  
Environment variables to set when starting a flow run.

1

2

3

Format

**Labels (Optional)**



# Run script to create the deployment with *.deploy()*

---

```
from prefect import flow

@flow(log_prints=True)
def my_flow(name: str = "World"):
    print(f"Hello {name}!")

if __name__ == "__main__":
    my_flow.from_source(
        source="https://github.com/biancaines/pal-2025-v1.git", # code stored in GitHub
        entrypoint="104/local-process-deploy-remote-code.py:my_flow",
    ).deploy(
        name="pal-local-process-deploy-remote-code",
        work_pool_name="pal-process-pool",
    )
```



## Start a worker

---

- In a new terminal window
- Watches for scheduled flow runs in the work pool

*prefect worker start --pool 'pal-process-pool'*



## Run the deployment using the CLI

---

Just like running a deployment with *.serve*

*prefect deployment run*  
*'my-flow/pal-local-process-deploy-remote-code'*





## See flow run logs in the UI or worker's terminal window

---

```
12:43:34.930 | INFO      | prefect.flow_runs.worker - Worker 'ProcessWorker c7a72edc-47  
56-4238-8083-bd615c763c60' submitting flow run '6e1140ff-7155-4288-8f8c-7d5ba0676c33'  
12:43:35.872 | INFO      | prefect.flow_runs.worker - Opening process...  
12:43:36.019 | INFO      | prefect.flow_runs.worker - Completed submission of flow run  
'6e1140ff-7155-4288-8f8c-7d5ba0676c33'
```



## At runtime:

---

1. Worker kicks off scheduled flow run
2. Pulls flow code from GitHub
3. Runs code in a local subprocess
4. Prefect monitors state
5. Subprocess exits



# Workers




# Workers

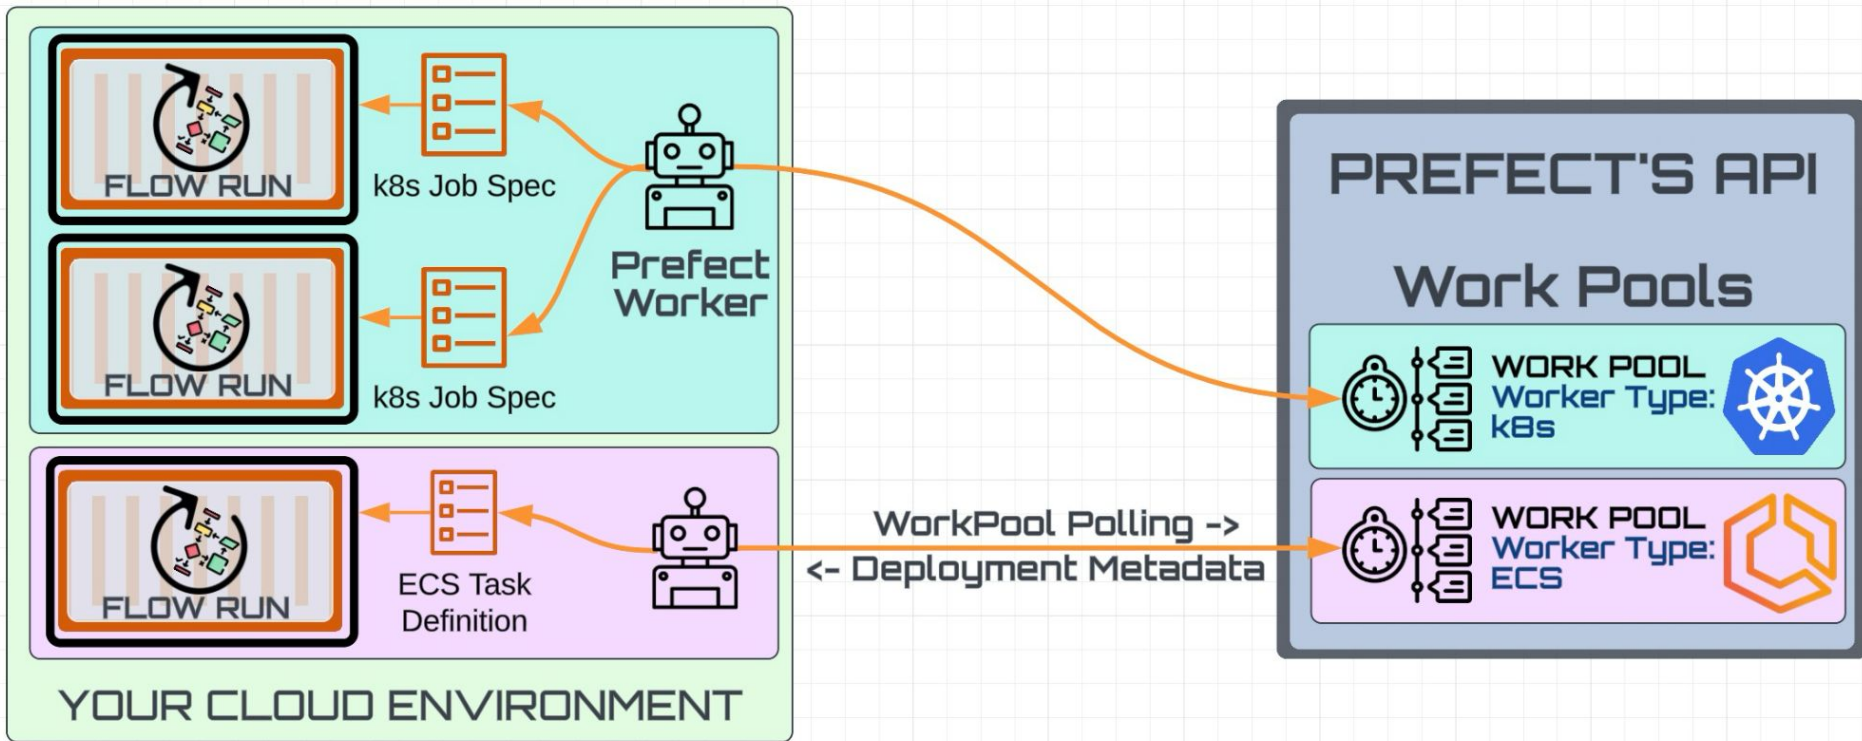
---

- Long-running process on the client
- Poll for scheduled flow runs from work pools
- Must match a work pool to pick up work
- If familiar with agents (old concept), workers are like smarter, typed agents



# WORKERS & WORK POOLS





# Example: Docker work pool & worker



# Why use Docker?

---

- Same operating environment everywhere
- Lighter weight than a VM
- Linux (generally)
- Portable
- Very popular
- All Prefect work pools other than Process use it





# Docker work pool

---

## Run a flow in a Docker container

1. Start Docker on your machine
2. Create a Docker type work pool
3. Start a worker that polls the work pool
4. Create a deployment that specifies the work pool
5. Run the deployment
6. Auto spins up a Docker container & runs flow in it



# Create a Docker work pool

## Hybrid

Hybrid work pools require workers to poll for and execute flow runs in your infrastructure.



### AWS Elastic Container Service

Execute flow runs within containers on AWS ECS. Works with EC2 and Fargate clusters. Requires an AWS account.



### Azure Container Instances

Execute flow runs within containers on Azure's Container Instances service. Requires an Azure account.



### Docker

Execute flow runs within Docker containers. Works well for managing flow execution environments via Docker images. Requires access to a running Docker daemon.



### Google Cloud Run

Execute flow runs within containers on Google Cloud Run. Requires a Google Cloud Platform account.



### Google Cloud Run V2

Execute flow runs within containers on Google Cloud Run (V2 API). Requires a Google Cloud Platform account.



### Google Vertex AI

Execute flow runs within containers on Google Vertex AI. Requires a Google Cloud Platform account.



### Kubernetes

Execute flow runs within jobs scheduled on a Kubernetes cluster. Requires a Kubernetes cluster.



# Docker work pool - base job template

**Base Job Template**

**Defaults**   Advanced

① The fields below control the default values for the base job template. These values can be overridden by deployments.

**Environment Variables (Optional)**  
Environment variables to set when starting a flow run.

1  
2  
3

Format

**Name (Optional)**  
Name given to infrastructure created by the worker using this job configuration.

**Image (Optional)**  
The image reference of a container image to use for created jobs. If not set, the latest Prefect image will be used.

docker.io/prefecthq/prefect:2-latest

**Labels (Optional)**  
Labels applied to infrastructure created by the worker using this job configuration.

1  
2  
3

Format



# Docker work pool - base job template

**Volumes (Optional)**  
A list of volume to mount into created containers.

1

2

3

/my/local/path:/path/in/container

Format

**Networks (Optional)**  
Docker networks that created containers should be connected to.

1

2

3

Format

**Memory Limit (Optional)**  
Memory limit of created containers. Accepts a value with a unit identifier (e.g. 100000b, 1000k, 128m, 1g.) If a value is given without a unit, bytes are assumed.

**Privileged (Optional)**  
Give extended privileges to created container.

☐

**Auto Remove (Optional)**  
If set, containers will be deleted on completion.

☒

**Network Mode (Optional)**  
The network mode for the created containers (e.g. host, bridge). If 'networks' is set, this cannot be set.

**Memory Swap Limit (Optional)**  
Total memory (memory + swap), -1 to disable swap. Should only be set if `mem_limit` is also set. If `mem_limit` is set, this defaults to allowing the container to use as much swap as memory. For example, if `mem_limit` is 300m and `memswap_limit` is not set, containers can use 600m in total of memory and swap.

**Stream Output (Optional)**  
If set, the output from created containers will be streamed to local standard output.

☒

**Image Pull Policy (Optional)**  
The image pull policy to use when pulling images.



# Package flow code into a Docker image with *.deploy()*

---

```
from prefect import flow

@flow(log_prints=True)
def buy():
    print("Buying securities")

if __name__ == "__main__":
    buy.deploy(
        name="my-code-in-an-image-deployment",
        work_pool_name="my-docker-pool",
        image="discdiver/local-image:1.0",
        push=False,
    )
```

! *.from\_source()* method not needed if baking flow code into image



## `.deploy()` method

---

Creates a Docker image with your flow code baked in by default!

- Specify the image name
- Specify `push=False` to not push image to registry
- Best practice: create a `requirements.txt` file with pinned package versions to install into image



## Docker type worker

---

Start a Docker type worker to connect to a work pool named *my-docker-pool*

```
prefect worker start -p my-docker-pool
```

If you want to make sure you have the packages needed:

```
prefect worker start -p my-docker-pool --install-policy always
```



# Docker

---

- Prefect provides base Docker images
- You can customize base image





# Docker

---

- Run your deployment
- Worker pulls image and spins up Docker container
- Flow code runs in Docker container and exits 🚀



# Docker

## See container in Docker Desktop if running locally

### Containers

[Give feedback](#) 

A container packages up code and its dependencies so the application runs quickly and reliably from one computing environment to another. [Learn more](#)

☐ Only show running containers

 Search



NAME

IMAGE

STATUS

PORT(S)

STARTED

ACTIONS



**nano-tortoise**

26fed9f8e268



[prefecthq/p](#) Exited



# Docker

---

## Reminders:

- Docker installed & **running**
- prefect-docker package installed
- Start a worker to poll for scheduled runs



# Hybrid work pool types

---

1. Process (local subprocess)
2. Docker
3. Serverless options such as ECS, ACI, GCR, VertexAI
4. Kubernetes

\* Worker required for all



# Push work pools



# Push work pools

---

## **Serverless. No worker required.**

- AWS ECS
- Google Cloud Run
- Azure Container Instances
- Modal
- Coiled



## Push work pools

---

Prefect will create everything for you with *--provision-infra*

Prerequisites:

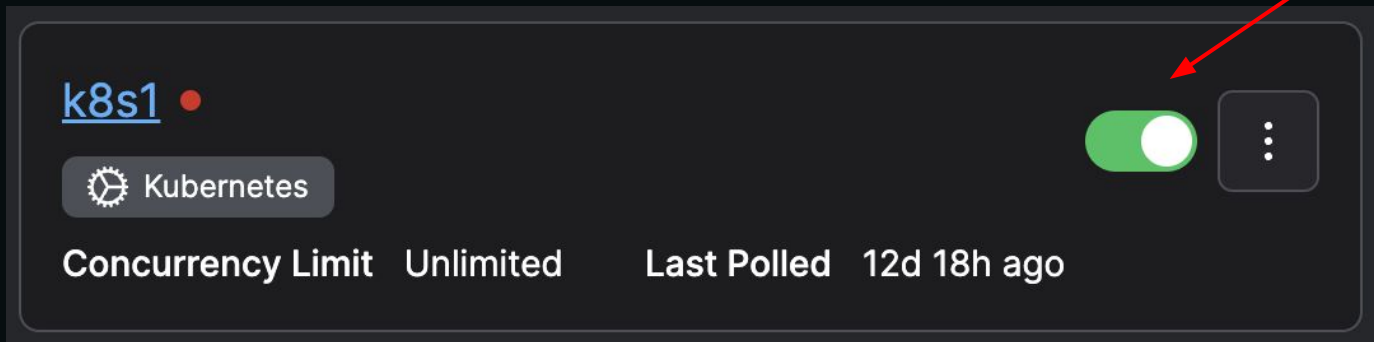
- Cloud provider account
- CLI tool installed
- Authenticated locally

*prefect work-pool create --type modal:push --provision-infra my-modal-pool*



# Pause **scheduled runs** for work pools from UI (or CLI)

---





## 104 Recap

---

You've seen how to

- Create work-pool based deployments! 🎉
- Use the hybrid model with workers
- Bake flow code into Docker images
- Run flows on a variety of infrastructure
- Pause and resume work pools




# Lab 104



# 104 Lab

---

- Let's make one of our weather forecast workflows more powerful
- Create a deployment with `.deploy()` that uses a **Prefect Managed** or **Process** work pool. Reminder, **Managed** is Prefect Cloud only
- Create work pool from the UI
  - Create a deployment that references flow code stored in your own GitHub repository
    - Use your earlier fetch weather flow if you like
    - **!** Push your code to your GitHub repo manually
  - If using a **Process** work pool start a worker to pick up scheduled flow runs
  - Run it! 



## 104 Lab Extensions

---

**Stretch 1:** Pause and resume the work pool from the UI.

**Stretch 2:** Experiment with adjusting fields in a work pool base job template.

**Stretch 3:** If you have Docker installed:

Create a deployment where you bake your flow code into a Docker image with `.deploy()`.

- Don't push the image (or log in + push to DockerHub).

Don't forget to:

- Start Docker on your machine
- Create a Docker work pool

